

SIMD Parallelism

- Vector architectures
- SIMD extensions
- Graphics Processor Units (GPUs)
- For x86 processors:
 - Expect two additional cores per chip per year
 - SIMD width to double every four years
 - Potential speedup from SIMD to be twice that from MIMD!

Chapter 4

Data-Level Parallelism in Vector, SIMD, and GPU Architectures



Vector Architectures

- Basic idea:
 - Read sets of data elements into “vector registers”
 - Operate on those registers
 - Disperse the results back into memory
- Registers are controlled by compiler
 - Used to hide memory latency
 - Leverage memory bandwidth

Vector Architectures

- Basic idea:
 - Read sets of data elements into “vector registers”
 - Operate on those registers
 - Disperse the results back into memory
- Registers are controlled by compiler
 - Used to hide memory latency
 - Leverage memory bandwidth

Introduction

- SIMD architectures can exploit significant data-level parallelism for:
 - Matrix-oriented scientific computing
 - Media-oriented image and sound processors
- SIMD is more energy efficient than MIMD
 - Only needs to fetch one instruction per data operation
 - Makes SIMD attractive for personal mobile devices
- SIMD allows programmer to continue to think sequentially

Vector Execution Time

- Execution time depends on three factors:
 - Length of operand vectors
 - Structural hazards
 - Data dependencies
- RV64V functional units consume one element per clock cycle
 - Execution time is approximately the vector length
- $Convey(编程时间)$
 - Set of vector instructions that could potentially execute together



VMIPS

- Example architecture: RV64V
 - Loosely based on Cray-1
 - 32 62-bit vector registers
 - Register file has 16 read ports and 8 write ports
 - Vector functional units
 - Fully pipelined
 - Data and control hazards are detected
 - Vector load-store unit
 - Fully pipelined
 - One word per clock cycle after initial latency
 - Scalar registers
 - 31 general-purpose registers
 - 32 floating-point registers



Chimes

- Sequences with read-after-write dependency hazards placed in same convey via *chaining*
- *Chaining*
 - Allows a vector operation to start as soon as the individual elements of its vector source operand become available
- *Chime*
 - Unit of time to execute one convey
 - m conveys executes in m chimes for vector length n
 - For vector length of n , requires $m \times n$ clock cycles



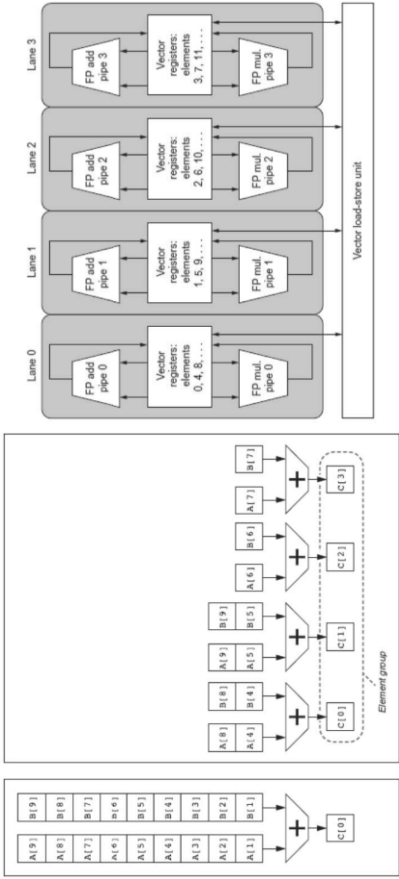
VMIPS Instructions

- `.vw`: two vector operands
- `.vs` and `.sv`: vector and scalar operands
- LV/SV: vector load and vector store from address
- Example: DAXPY
 - `vsetdcfg 4*FP64` # Enable 4 DP FP vregs
 - `fld f0,a` # Load scalar a
 - `vld v0,x5` # Load vector X
 - `vmul v1,v0,f0` # Vector-scalar mult
 - `vld v2,x6` # Load vector Y
 - `vadd v3,v1,v2` # Vector-vector add
 - `vst v3,x6` # Store the sum
 - `vdisable` # Disable vector regs
- 8 instructions, 258 for RV64V (scalar code)



Multiple Lanes

- Element n of vector register A is “hardwired” to element n of vector register B
- Allows for multiple hardware lanes



Example

```
vld    v0,x5      # Load vector X
vmul   v1,v0,f0    # Vector-scalar multiply
vld    v2,x6      # Load vector Y
vadd   v3,v1,v2    # Vector-vector add
vst    v3,x6      # Store the sum
```

Convoys:

```
1      vld          vmul
2      vld          vadd
3      vst
```

3 chimes, 2 FP ops per result, cycles per FLOP = 1.5
For 64 element vectors, requires $32 \times 3 = 96$ clock cycles

Vector Length Register

```
for (i=0; i < n; i=i+1) Y[i] = a * X[i] + Y[i];
vsetdcfg 2 DP FP # Enable 2 64b F1,Pt. registers
fid f0,a      # Load scalar a
loop: setv1 t0,a0
vld v0,x5     # Load vector X
slli t1,t0,3  # t1 = v1 * 8 (in bytes)
add x5,x5,t1  # Increment pointer to X by v1*8
vmul v0,v0,f0 # Vector-scalar mult
vld v1,x6     # Load vector Y
vadd v1,v0,v1 # Vector-vector add
sub a0,a0,t0  # n -= v1 (t0)
vst v1,x6     # Store the sum into Y
add x6,x6,t1  # Increment pointer to Y by v1*8
bnez a0,loop  # Repeat if n != 0
vdisable     # Disable vector regs
```

Challenges

- Start up time
 - Latency of vector functional unit
 - Assume the same as Cray-1
 - Floating-point add => 6 clock cycles
 - Floating-point multiply => 7 clock cycles
 - Floating-point divide => 20 clock cycles
 - Vector load => 12 clock cycles
- Improvements:
 - > 1 element per clock cycle
 - Non-64 wide vectors
 - IF statements in vector code
 - Memory system optimizations to support vector processors
 - Multiple dimensional matrices
 - Sparse matrices
 - Programming a vector computer

Stride

- Consider:
for (i = 0; i < 100; i=i+1)
 for (j = 0; j < 100; j=j+1) {
 A[i][j] = 0.0;
 for (k = 0; k < 100; k=k+1)
 A[i][j] = A[i][j] + B[i][k] * D[k][j];
 }

Must vectorize multiplication of rows of B with columns of D

Use *non-unit stride*

Bank conflict (stall) occurs when the same bank is hit faster than bank busy time:
 - #banks / LCM(stride,#banks) < bank busy time



Vector Mask Registers

- Consider:
for (i = 0; i < 64; i=i+1)
 if (X[i] != 0)
 X[i] = X[i] - Y[i];

Use predicate register to “disable” elements:
vsetdcfg 2*FP64 # Enable 2 64b FP vector regs
vsetpcfgi 1 # Enable 1 predicate register
vld v0,x5 # Load vector X into v0
vld v1,x6 # Load vector Y into v1
fmv.d.x f0,x0 # Put (FP) zero into f0
vpne p0,v0,f0 # Set p0(i) to 1 if v0(i)=f0
vsub v0,v0,v1 # Subtract under vector mask
vst v0,x5 # Store the result in X
vdisable # Disable vector registers
vpdisable # Disable predicate registers



Scatter-Gather

- Consider:
for (i = 0; i < n; i=i+1)
 A[K[i]] = A[K[i]] + C[M[i]];

Use index vector:
vsetdcfg 4*FP64 # 4 64b FP vector registers
vld v0,x7 # Load K[]
vldx v1,x5,v0 # Load A[K[]]
vld v2,x28 # Load M[]
vldi v3,x6,v2 # Load C[M[]]
vadd v1,v1,v3 # Add them
vstx v1,x5,v0 # Store A[K[]]
vdisable # Disable vector registers



Memory Banks

- Memory system must be designed to support high bandwidth for vector loads and stores
- Spread accesses across multiple banks
 - Control bank addresses independently
 - Load or store non sequential words (need independent bank addressing)
 - Support multiple vector processors sharing the same memory
- Example:
 - 32 processors, each generating 4 loads and 2 stores/cycle
 - Processor cycle time is 2.167 ns, SRAM cycle time is 15 ns
 - How many memory banks needed?
 - $32 \times (4+2) \times 15 / 2.167 = \sim 1330$ banks



- Implementations:
 - Intel MMX (1996)
 - Eight 8-bit integer ops or four 16-bit integer ops
 - Streaming SIMD Extensions (SSE) (1999)
 - Eight 16-bit integer ops
 - Four 32-bit integer/fp ops or two 64-bit integer/fp ops
 - Advanced Vector Extensions (2010)
 - Four 64-bit integer/fp ops
 - AVX-512 (2017)
 - Eight 64-bit integer/fp ops
 - Operands must be consecutive and aligned memory locations

- Compilers can provide feedback to programmers
- Programmers can provide hints to compiler

Benchmark name	Operations executed in vector mode, compiler-optimized	Operations executed in vector mode, with programmer aid	Speedup from hint optimization
BDNA	96.1%	97.2%	1.52
MG3D	95.1%	94.5%	1.00
FLO2	91.5%	88.7%	N/A
ARC3D	91.1%	92.0%	1.01
SPEC77	90.3%	90.4%	1.07
MDG	87.7%	94.2%	1.49
TRFD	69.8%	73.7%	1.67
DYFESM	68.8%	65.6%	N/A
ADM	42.9%	59.6%	3.60
OCEAN	42.8%	91.2%	3.92
TRACK	14.4%	54.6%	2.52
SPIE	11.5%	79.9%	4.06
QCD	4.2%	75.1%	2.15

- Example DXPY:

```
fld      f0,a          # Load scalar a
splat.4D f0,f0         # Make 4 copies of a
addi     x28,x5,#256    # Last address to load
Loop: fld.4D f1,0(x5)   # Load X[i] ... X[i+3]
fmul.4D  f1,f1,f0       # a x X[i] ... a x X[i+3]
fld.4D   f2,0(x6)       # Load Y[i] ... Y[i+3]
fadd.4D  f2,f2,f1       # a x X[i]+Y[i]...
          f2,f2,f1       # a x X[i+3]+Y[i+3]
fsd.4D   f2,0(x6)       # Store Y[i]... Y[i+3]
addi     x5,x5,#32      # Increment index to X
addi     x6,x6,#32      # Increment index to Y
bne      x28,x5,Loop    # Check if done
```


- Media applications operate on data types narrower than the native word size
 - Example: disconnect carry chains to “partition” adder
- Limitations, compared to vector instructions:
 - Number of data operands encoded into op code
 - No sophisticated addressing modes (strided, scatter-gather)
 - No mask registers

Graphical Processing Units

- Basic idea:
 - Heterogeneous execution model
 - CPU is the *host*, GPU is the *device*
 - Develop a C-like programming language for GPU
 - Unify all forms of GPU parallelism as *CUDA thread*
 - Programming model is “Single Instruction Multiple Thread”

Roofline Performance Model

- Basic idea:
 - Plot peak floating-point throughput as a function of arithmetic intensity
 - Ties together floating-point performance and memory performance for a target machine
 - Arithmetic intensity
 - Floating-point operations per byte read




Threads and Blocks

- A thread is associated with each data element
- Threads are organized into blocks
- Blocks are organized into a grid
- GPU hardware handles thread management, not applications or OS

Examples

- Attainable GFLOPs/sec = (Peak Memory BW × Arithmetic Intensity, Peak Floating Point Perf.)



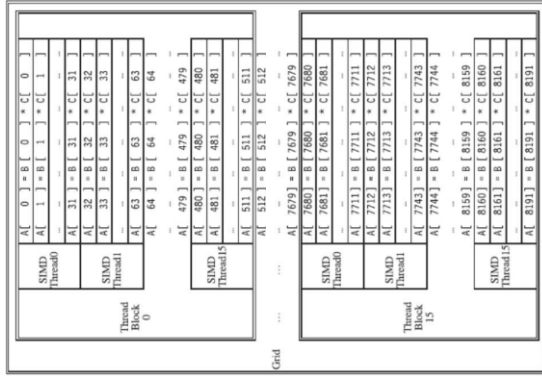
Terminology

- Each thread is limited to 64 registers
- Groups of 32 threads combined into a SIMD thread or “warp”
 - Mapped to 16 physical lanes
- Up to 32 warps are scheduled on a single SIMD processor
 - Each warp has its own PC
 - Thread scheduler uses scoreboard to dispatch warps
 - By definition, no data dependencies between warps
 - Dispatch warps into pipeline, hide memory latency
- Thread block scheduler schedules blocks to SIMD processors
 - Within each SIMD processor:
 - 32 SIMD lanes
 - Wide and shallow compared to vector processors

NVIDIA GPU Architecture

- Similarities to vector machines:
 - Works well with data-level parallel problems
 - Scatter-gather transfers
 - Mask registers
 - Large register files
- Differences:
 - No scalar processor
 - Uses multithreading to hide memory latency
 - Has many functional units, as opposed to a few deeply pipelined units like a vector processor

Example




Example

- Code that works over all elements is the grid
- Thread blocks break this down into manageable sizes
 - 512 threads per block
- SIMD instruction executes 32 elements at a time
- Thus grid size = 16 blocks
- Block is analogous to a strip-mined vector loop with vector length of 32
- Block is assigned to a multithreaded SIMD processor by the thread block scheduler
- Current-generation GPUs have 7-15 multithreaded SIMD processors

Conditional Branching

- Like vector architectures, GPU branch hardware uses internal masks
- Also uses
 - Branch synchronization stack
 - Entries consist of masks for each SIMD lane
 - I.e. which threads commit their results (all threads execute)
 - Instruction markers to manage when a branch diverges into multiple execution paths
 - Push on divergent branch
 - ...and when paths converge
 - Act as barriers
 - Pops stack
- Per-thread-lane 1-bit predicate register, specified by programmer

GPU Organization



NVIDIA GPU Memory Structures

- Each SIMD Lane has private section of off-chip DRAM
 - "Private memory"
 - Contains stack frame, spilling registers, and private variables
- Each multithreaded SIMD processor also has local memory
 - Shared by SIMD lanes / threads within a block
- Memory shared by SIMD processors is GPU Memory
 - Host can read and write GPU memory

NVIDIA Instruction Set Arch.

- ISA is an abstraction of the hardware instruction set
 - "Parallel Thread Execution (PTX)"
 - opcode.type d,a,b,c;
 - Uses virtual registers
 - Translation to machine code is performed in software
 - Example:


```
shl.s32    R8, blockIdx, 9      ; Thread Block ID * Block size (512 or 2^9)
add.s32    R8, R8, threadIdx   ; R8 = i = my CUDA thread ID
ld.global.f64 RD0, [X+R8]      ; RD0 = X[i]
ld.global.f64 RD2, [Y+R8]      ; RD2 = Y[i]
mul.f64 RD0, RD0, RD4          ; Product in RD0 = RD0 * RD4 (scalar a)
add.f64 RD0, RD0, RD2          ; Sum in RD0 = RD0 + RD2 (Y[i])
st.global.f64 [Y+R8], RD0      ; Y[i] = sum (X[i]*a + Y[i])
```


Vector Architectures vs GPUs

- SIMD processor analogous to vector processor, both have MIMD
- Registers
 - RV64V register file holds entire vectors
 - GPU distributes vectors across the registers of SIMD lanes
 - RV64 has 32 vector registers of 32 elements (1024)
 - GPU has 256 registers with 32 elements each (8K)
 - RV64 has 2 to 8 lanes with vector length of 32, chime is 4 to 16 cycles
- SIMD processor chime is 2 to 4 cycles
- GPU vectorized loop is grid
- All GPU loads are gather instructions and all GPU stores are scatter instructions



Pascal Architecture Innovations

- Each SIMD processor has
 - Two or four SIMD thread schedulers, two instruction dispatch units
 - 16 SIMD lanes (SIMD width=32, chime=2 cycles), 16 load-store units, 4 special function units
 - Two threads of SIMD instructions are scheduled every two clock cycles
- Fast single-, double-, and half-precision
- High Bandwidth Memory 2 (HBM2) at 732 GB/s
- NVLink between multiple GPUs (20 GB/s in each direction)
- Unified virtual memory and paging support




SIMD Architectures vs GPUs

- GPUs have more SIMD lanes
- GPUs have hardware support for more threads
- Both have 2:1 ratio between double- and single-precision performance
- Both have 64-bit addresses, but GPUs have smaller memory
- SIMD architectures have no scatter-gather support



Pascal Multithreaded SIMD Proc.



Fallacies and Pitfalls

- GPUs suffer from being coprocessors
 - GPUs have flexibility to change ISA
- Concentrating on peak performance in vector architectures and ignoring start-up overhead
 - Overheads require long vector lengths to achieve speedup
- Increasing vector performance without comparable increases in scalar performance
- You can get good vector performance without providing memory bandwidth
- On GPUs, just add more threads if you don't have enough memory performance